

# Lecture 3 Additional Slides

CSE 344, Winter 2014

Sudeepa Roy

- Note:

These slides mostly contain the same material as the lecture notes, and were used as a substitute for a whiteboard in class.

- Also please go over all the sql commands and the comments in lecture02 and lecture03 notes and let us know if any of the queries or explanations is not clear.

- A few very nice observations from you in class!
  - Check if you using the right quote ' and not an automatic correction by an editor '
  - Check if the font size is proper in sqlite window or whether it is truncating any field: “photography” vs. “photograph”

# Ex. In SQLite

## Product

pname	price	category	manufacturer
Gizmo	19.99	gadget	GizmoWorks
PowerGizmo	29.99	gadget	GizmoWorks
SingleTouch	149.99	photography	Canon
MultiTouch	199.99	photography	Hitachi
SuperGizmo	49.99	gadget	Hitachi

## Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan
Apple	USA

**manufacturer**  
references **cname**

- Product(pname, price, category, manufacturer)
- Company(cname, country)

# Selection queries

- Product(pname, price, category, manufacturer)
- Company(cname, country)
- Select a subset of rows
- Condition specified by WHERE clause
- Ex:
  - select \*
  - From Product
  - where price > 100.0;
- Ex:
  - select \*
  - From Product
  - Where pname like '%e%';

# Projection queries

- Product(pname, price, category, manufacturer)
  - Company(cname, country)
- 
- Keep a subset of the attributes/columns
  - Attributes specified by SELECT clause
  - Ex:
    - select price, category
    - from Product;

# DISTINCT

- Product(pname, price, category, manufacturer)
- Company(cname, country)

- Duplicates:
  - select category
  - from Product;
- Eliminates duplicates
  - select distinct category
  - from Product;
- Once again, set vs. bag

# Order By

- Product(pname, price, category, manufacturer)
- Company(cname, country)

- order alphabetically by name:
- order by price descending
- order by manufacturer, then price descending



# BASIC SQL Query Evaluation

```
SELECT <attr>  
FROM <reln>  
WHERE <condn>
```

```
(optional)  
ORDER BY <attr2>  
        <asc/desc>
```

## Sequence in evaluation

1. **FROM:** for each tuple in <reln>
2. **WHERE:** apply <condn>
3. **SELECT:** <attr>

# ORDER BY and DISTINCT - 1

```
SELECT <attr>  
FROM <reln>  
WHERE <condn>
```

```
(optional)  
ORDER BY <attr2>  
        <asc/desc>
```

## Sequence in evaluation

1. **FROM**: for each tuple in <reln>
2. **WHERE**: apply <condn>
3. **ORDER BY**: <attr2>
4. **SELECT**: <attr>

# ORDER BY and DISTINCT - 2

- What happens if we order on an attribute that we do NOT return ?
- First, let's try:
  - `select * from Product order by manufacturer;`
- Now, let's try:
  - `select category from Product order by manufacturer;`
- What happens if we also do DISTINCT ?
  - `select distinct category from Product order by manufacturer;`
- In SQL, all attributes in ORDER BY must appear in SELECT if DISTINCT is used, should have been an error. Sqlite does not enforce this (another alert).

# JOINS

- Product(pname, price, category, manufacturer)
- Company(cname, country)

- What should the following query return?
  - select pname, price
  - from Product P, Company C
  - where P.**manufacturer**=C.**cname** and country='Japan' and price < 150;

Join predicate



- Your answer in class 😊
- Single touch, 149.99
- Supergizmo, 49.99

Ex. 1: Retrieve all American company names that manufacture products in the 'gadget' category

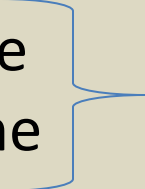
- Product(pname, price, category, manufacturer)
- Company(cname, country)
- **Your answer in class 😊**

```
SELECT distinct P.manufacturer
FROM Product P, Company C
WHERE P.manufacturer = C.cname
And C.country = 'USA'
And P.category = 'gadget'
```

## Ex. 2: Retrieve all Japanese company names that manufacture products in both the 'gadget' and the photography category

- Product(pname, price, category, manufacturer)
- Company(cname, country)
- **Your answer in class** 😊

```
SELECT distinct cname
FROM Product P1,
Company, Product P2
WHERE country = 'Japan' AND
P1.category = 'gadget' AND
P2.category = 'photography'
And P1.manufacturer = cname
AND P2.manufacturer = cname
```



Note:

1. The third condition  
P1.manufacturer =  
P2.manufacturer is not needed
2. We could replace the last  
condition by P1.manufacturer =  
P2.manufacturer
3. i.e. only two equality checks are  
needed and not three
4. Why? See next 4 slides and  
think!

# Join as a cartesian product followed by selection and projection (NEW SLIDE-1)

- You should think this way when writing complicated SQL queries or finding answers to a given query.
- More on this when we learn Relational Algebra (Lecture 9)
- Consider this example.

```
select R.a  
from R, S, T  
where R.a=S.a AND S.a <> b
```

R	S	T
A	A	B
1	2	1
2	3	2
3	4	4

# Join as a cartesian product followed by selection and projection (NEW SLIDE-2)

R	S	T
A	A	B
1	2	3
2	3	

```
select R.a  
from R, S, T  
where R.a=S.a AND S.a <> b
```

R.A	S.A	T.B
1	2	3
1	3	3
2	2	3
2	3	3

$R \times S \times T$

- STEP 1: The cartesian product of these tables will be computed (see the nested loop semantics later)
- NOTE: DBMSs will almost never evaluate queries in this inefficient way. More on this when we learn query plans



# Join as a cartesian product followed by selection and projection (NEW SLIDE-3)

R	S	T
A	A	B
1	2	3
2	3	

select R.a  
from R, S, T  
where R.a=S.a AND S.a <> b

- STEP 2: Apply condition in WHERE clause
- Only one tuple satisfies the condition (highlighted)

R.A	S.A	T.B
1	2	3
1	3	3
<b>2</b>	<b>2</b>	<b>3</b>
2	3	3

$R \times S \times T$

# Join as a cartesian product followed by selection and projection (NEW SLIDE-4)

R	S	T
A	A	B
1	2	3
2	3	

select R.a  
from R, S, T  
where R.a=S.a AND S.a <> b

R.A	S.A	T.B
1	2	3
1	3	3
<b>2</b>	<b>2</b>	<b>3</b>
2	3	3

$R \times S \times T$

- STEP 3: Now project on to the attributes in SELECT clause
- Final answer:  
2

# Now think!

- why the other two answers for Ex 2 that we considered in class did not work
  - (Japanese companies for both gadgets and photography)
- Option 1: country = 'Japan' and (category = 'gadget' OR category = 'photography')
  - Ans: If any of these two categories exists in the cartesian product with country = Japan, it will be returned
- Option 2: country = 'Japan' and (category = 'gadget' AND category = 'photography')
  - Ans: Category for any tuple in the cartesian product cannot be both gadget and photography

# Joins may introduce duplicates

- Try:
  - select country
  - from Product, Company
  - where manufacturer=cname and category='gadget';
- Easy fix: USE DISTINCT
  - select distinct country
  - from Product, Company
  - where manufacturer=cname and category='gadget';

## Ex. 3: Find all countries that manufacture both a product under \$25 and a product over \$25 (Aliases)

- Product(pname, price, category, manufacturer)
- Company(cname, country)
- Answer :
- (First try yourself and then see the answer in the notes.)

# JOINS: Nested Loop Semantics for SQL

- Query:
  - SELECT a1, a2, ..., ak
  - FROM R1 AS x1, R2 AS x2, ..., Rm AS xm
  - WHERE Cond

Although the quer  
processor will  
ALMOST NEVER  
evaluate the query  
this way!

- Semantics:
  - for a1 in R1 do
  - for a2 in R2 do
  - for a3 in R3 do –
  - ...
  - for an in Rm do
  - if Cond(a1, ...ak) is true
  - then output(a1,...,ak)

FROM

WHERE

SELECT

# What does this query compute?

R	S	T
A	A	A
1	2	1
2	3	2
3	4	4

```
select distinct R.a  
from R, S  
where R.a=S.a;
```

# What does this query compute?

R	S	T
A	A	A
1	2	1
2	3	2
3	4	4

ANS: R intersects S

```
select distinct R.a  
from R, S  
where R.a=S.a;
```



# What does this query compute?

R	S	T
A	A	A
1	2	1
2	3	2
3	4	4

select distinct T.a  
from R, S, T  
where R.a=T.a or S.a = T.a

# What does this query compute?

R	S	T
A	A	A
1		1
2		2
3		4

- you might think it is:  $(R \cup S) \cap T$
- but think again!
- what happens if say  $S = \emptyset$ , i.e. no tuples in S
- The result should have been  $(R \cap T)$
- But we get empty set.
- Nested loop semantics explains this!

select distinct T.a  
from R, S, T  
where R.a=T.a or S.a = T.a

Ans: the query returns  
 $(R \cup S) \cap T$   
if R,S are non-empty.  
otherwise it returns the empty set

# NULL In SQLite

## Product

pname	price	category	manufacturer
Gizmo	19.99	gadget	GizmoWorks
PowerGizmo	29.99	gadget	GizmoWorks
SingleTouch	149.99	photography	Canon
MultiTouch	199.99	photography	Hitachi
SuperGizmo	49.99	gadget	Hitachi
iPad 5	NULL	gadget	Apple

## Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan
Apple	USA

**SELECT gadegts with price  
< 25 and >=25  
Ipad is nowhere!**

- Product(pname, price, category, manufacturer)
- Company(cname, country)

# Conditions involving NULL

- We need to evaluate in SQL conditions like this:
- $(\text{price} < 25) \text{ and } (\text{category} = \text{'gadget'}) \text{ or } (\text{manufacturer} = \text{'Apple'})$
- Suppose  $\text{price} = 19$ ,  $\text{category} = \text{NULL}$ , and  $\text{manufacturer} = \text{NULL}$
- Is the predicate true or false?

# 3-valued logic

- FALSE = 0 E.g.  $\text{price} < 25$  is FALSE when  $\text{price} = 99$
- UNKNOWN = 0.5 E.g.  $\text{price} < 25$  is UNKNOWN when  $\text{price} = \text{NULL}$
- TRUE = 1 E.g.  $\text{price} < 25$  is TRUE when  $\text{price} = 19$
  
- $C1 \text{ AND } C2$  means  $\min(C1, C2)$
- $C1 \text{ OR } C2$  means  $\max(C1, C2)$
- not C means  $1 - C$

# Compute the truth value

- $(\text{price} < 25) \text{ and } (\text{category} = \text{'gadget'}) \text{ or } (\text{manufacturer} = \text{'Apple'})$
- Suppose  $\text{price} = 19$ ,  $\text{category} = \text{NULL}$ , and  $\text{manufacturer} = \text{NULL}$

- Answer:

# Output of a query for NULL

- The rule for SELECT ... FROM ... WHERE C is the following:
  - if C = TRUE then include the row in the output
  - if C = FALSE or C = unknown then do not include it

# Outer JOIN

- JOIN = INNER JOIN
- LEFT OUTER JOIN:
  - include everything on the left, fill in the right part with NULL values
- Similarly (FULL)/RIGHT OUTER JOIN

More on outer join in Lec 4 (Monday)

pname	price	category	manufacturer
Gizmo	19.99	gadget	GizmoWorks
PowerGizmo	29.99	gadget	GizmoWorks
SingleTouch	149.99	photography	Canon
MultiTouch	199.99	photography	Hitachi
SuperGizmo	49.99	gadget	Hitachi
iPad 5	NULL	gadget	Apple

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan
Apple	USA
<b>Google</b>	<b>USA</b>

Product

Company