

Introduction to Data Management

CSE 344

Lectures 5: Aggregates in SQL

Daniel Halperin

Announcements

- Webquiz 2 posted this morning
- Homework 1 is due on Thursday (01/16)

(Random detour:) Who is this?



http://content.lib.washington.edu/cdm4/item_viewer.php?CISOROOT=/portraits&CISOPTR=117&CISOBX=1&REC=5

Does this help?



http://content.lib.washington.edu/cdm4/item_viewer.php?CISOROOT=/uwcampus&CISOPTR=1649

Winlock W Miller (of Miller Hall :)

- UW Regent (managers of univ.) for 35 years between 1913 and 1953
 - Usually full of executives from major institutions
 - Current Board of Regents Chair is former Alaska Airlines CEO, etc.
- Winlock, WA is named after him
- Father was Gen'l William Winlock Miller (confusing, I know), first mayor of Olympia and land speculator.
- (I think) WA has some interesting history!

Refresh your memory

```
> SELECT * FROM Purchase;
```

pid	product	price	quantity	month
-----	-----	-----	-----	-----
1	bagel	1.99	20	september
2	bagel	2.5	12	december
3	banana	0.99	9	september
4	banana	1.59	9	february
5	gizmo	99.99	5	february
6	gizmo	99.99	3	march
7	gizmo	49.99	3	april
8	gadget	89.99	3	january
9	gadget	89.99	3	february
10	gadget	49.99	3	march
11	orange	NULL	5	may
12	orange	1.29	34	january

Refresh your memory

How do we...

- Compute the total number of sales?
- Compute the total number of products sold?
- Compute the total number of each product sold?
- Compute the gross \$ spent on of each product?
(qty * price)
- Compute the average gross \$ of each product?
(2 ways)

Refresh your memory

How do we...

- Compute the gross monthly sales in \$?
(units * price/unit)
- Sort the months from most sales to least?
- Find all the unique prices? (2 ways)

HAVING Clause

Same query as earlier, except that we consider only products that had at least 30 sales.

```
SELECT    product, sum(price*quantity)
FROM      Purchase
WHERE     price > 1
GROUP BY  product
HAVING    Sum(quantity) > 30
```

HAVING clause contains conditions on aggregates.

WHERE vs HAVING

- WHERE condition is applied to individual rows
 - The rows may or may not contribute to the aggregate
 - No aggregates allowed here
- HAVING condition is applied to the entire group
 - Entire group is returned, or not at all
 - May use aggregate functions in the group

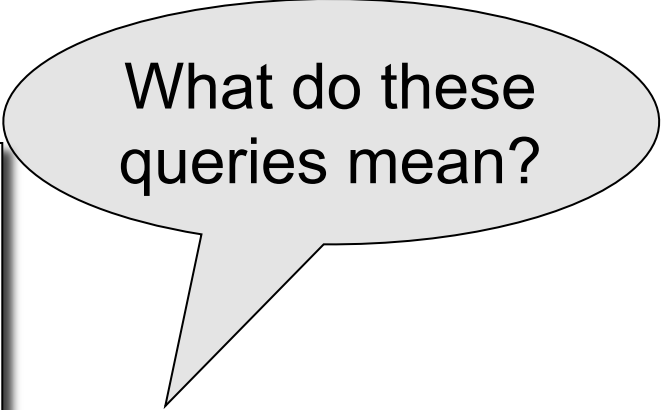
Aggregates and Joins

```
create table Product
(pid int primary key,
 pname varchar(15),
 manufacturer varchar(15));

insert into product values(1, 'bagel', 'Sunshine Co. ');
insert into product values(2, 'banana', 'BusyHands ');
insert into product values(3, 'gizmo', 'GizmoWorks ');
insert into product values(4, 'gadget', 'BusyHands ');
insert into product values(5, 'powerGizmo', 'PowerWorks ');
```

Aggregate + Join Example

```
SELECT x.manufacturer, count(*)  
FROM Product x, Purchase y  
WHERE x.pname = y.product  
GROUP BY x.manufacturer
```



What do these queries mean?

```
SELECT x.manufacturer, y.month, count(*)  
FROM Product x, Purchase y  
WHERE x.pname = y.product  
GROUP BY x.manufacturer, y.month
```

General form of Grouping and Aggregation

SELECT	S
FROM	R_1, \dots, R_n
WHERE	C1
GROUP BY	a_1, \dots, a_k
HAVING	C2



Why ?

S = may contain attributes a_1, \dots, a_k and/or any aggregates but NO OTHER ATTRIBUTES

C1 = is any condition on the attributes in R_1, \dots, R_n

C2 = is any condition on aggregate expressions and on attributes a_1, \dots, a_k

Semantics of SQL With Group-By

SELECT	S
FROM	R_1, \dots, R_n
WHERE	C1
GROUP BY	a_1, \dots, a_k
HAVING	C2

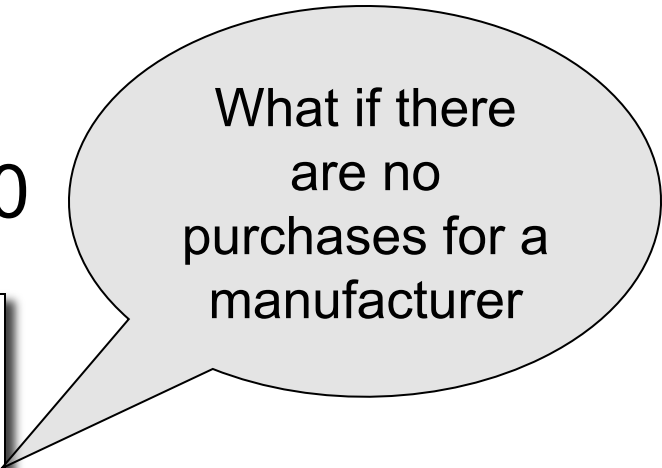
Evaluation steps:

1. Evaluate FROM-WHERE using Nested Loop Semantics
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C2 to each group (may have aggregates)
4. Compute aggregates in S and return the result

Empty Groups

- In the result of a group by query, there is one row per group in the result
- No group can be empty!
- In particular, count(*) is never 0

```
SELECT x.manufacturer, count(*)  
FROM Product x, Purchase y  
WHERE x.pname = y.product  
GROUP BY x.manufacturer
```



What if there
are no
purchases for a
manufacturer

Empty Groups: Example

```
SELECT product, count(*)  
FROM purchase  
GROUP BY product
```

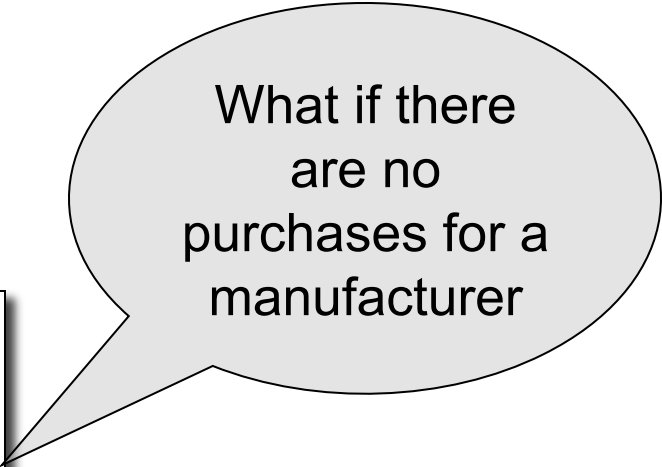
5 groups in our
example dataset

```
SELECT product, count(*)  
FROM purchase  
WHERE price > 2.0  
GROUP BY product
```

3 groups in our
example dataset

Empty Group Problem

```
SELECT x.manufacturer, count(*)  
FROM Product x, Purchase y  
WHERE x.pname = y.product  
GROUP BY x.manufacturer
```



What if there
are no
purchases for a
manufacturer

Empty Group Solution: Outer Join

```
SELECT x.manufacturer, count(y.pid)
FROM Product x LEFT OUTER JOIN Purchase y
ON x.pname = y.product
GROUP BY x.manufacturer
```

- 1) List all manufacturers with more than 10 items sold. Return the manufacturer name and the number of items sold.