

CSE 344 Section 3 Worksheet

More Nested Queries

1) Consider a photo sharing Website, where users can post pictures, as well as comment and rate other user's pictures. The schema is:

Users(uid, name)
Comment(uid, pid, score, txt)
Picture(pid, author, img)

The database has the following constraints:

- ... Comment.uid is a foreign key to Users
- ... Comment.pid is a foreign key to Picture
- ... Picture.author is a foreign key to Users
- ... Comment.score is an integer number between 1 and 10
- ... All attributes are NOT NULL

A picture is considered highly rated if it received at least one score of 10, from a user other than its author. A cautious user is a user who commented only on highly rated pictures. (A user who did not comment at all is also cautious.) Write a SQL query that finds all cautious users. Your query should return a list of uid, name pairs

```
select *  
from users  
where uid not in  
    (select y.uid -- these are the non-cautious users  
    from comment y  
    where y.pid not in  
        (select u.pid -- the pictures that are highly rated  
        from comment u, picture v  
        where u.pid = v.pid and u.score = 10 and u.pid != v.author))
```

A negation can be avoided by using an aggregate, e.g. having(max(score) < 10.)

2) (SAME SCHEMA AS LAST WEEK) A database with the following schema stores a collection of Webpages and the words they contain, and a collection of dictionaries in several languages and the words in those languages:

Occurs(url, word)
Dictionary(language, word)

- url represents a Webpage.
- Every Webpage may contain several words, and every word may occur in several Webpages.
- Every language may contain several words, and every word may occur in several languages
- There are no nulls in the database.

Write a SQL query that computes, for each Webpage, the language that has the largest number of word on that page. For example, if a page has 100 words in French and 50 words in English (these two sets of words may be overlapping), then your query will return French for that page. If a Webpage has only words that do not occur in any language at all, then you do not need to return that Webpage.

```
select distinct z.url,  
      (select y.language  
       from Occur x, Dictionary y  
       where x.word = y.word and x.url = z.url  
       group by y.language  
       order by count(*) desc  
       limit 1)  
from Occurs z
```

3) Consider the following social network database:

Person(pid, name)

Relationship(pid1, pid2, type)

Where:

- Person.pid is a key.
- Relationship.pid1 and Relationship.pid2 are foreign keys to Person.
- Relationship.type is either 'friend' or 'enemy'.

Keep in mind that Relationship is not symmetric: if p1 is a friend of p2, that does not mean p2 is a friend of p1. It is not transitive either: if p1 is a friend of p2 who is a friend of p3, it doesn't mean p1 is a friend of p3.

(a) (12 points) A second degree friend is a friend's friend. Write a SQL query that computes for each person the total number of their second degree friends. Your query should return answers of the form: pid, name, count. Cryptic hint: "not every person has friends, but you have to count everyone's second degree friends".

```
select x.pid, x.name, count(distinct z.pid2)
from Person x left outer join Relationship y
    on x.pid = y.pid1 and y.type = 'friend'
    left outer join Relationship z
    on y.pid2 = z.pid1 and z.pid2 != x.pid
    and z.type = 'friend'
group by x.pid, x.name
```

(b) (12 points) Write a SQL query that returns all persons who have at least 12 common friends with "Mary". Your query should return answers of the form: pid1, pid2, name, where pid1 is Mary's pid and pid2 is that of a person who has 12 or more common friends (meaning there are at least 12 persons p such that pid1 and p are friends, and pid2 and p are friends). If there are multiple people called Mary, then you will report each of them.

```
select x.pid, y.pid, y.name  
from Person x, Person y, Relationship u, Relationship v  
where x.pid = u.pid1 and y.pid = v.pid1 and u.pid2 = v.pid2  
      and u.type = 'friend' and v.type = 'friend'  
      and x.name = 'Mary'  
group by x.pid, y.pid, y.name  
having count(*) >= 12
```

(c) (12 points) Fred says: "my enemies' enemies are my friends". Prove that Fred is wrong: write a query that returns all Fred's enemies' enemies that are not his friends. Your query should return answers of the form: pid1, pid2, where pid1 is Fred's pid and pid2 represents an enemy's enemy that is not Fred's friend.

```
select x.pid, z.pid2  
from Person x, Relationship y, Relationship z  
where x.name = 'Fred'  
      and x.pid = y.pid1 and y.pid2 = z.pid1  
      and y.type = 'enemy' and z.type = 'enemy'  
      and z.pid2 not in (select u.pid2  
                          from Relationship u where u.type = 'friend'  
                          and u.pid1 = y.pid1)
```